



清华大学东南亚中心
Tsinghua Southeast Asia Center



清华大学深圳国际研究生院
Tsinghua Shenzhen International Graduate School

Introduction to GenAI

Happy Digital X

Happy Digital X | Tsinghua University

Today's Agenda

1 Data Governance

Privacy regulations and data management

2 Product Development

Lifecycle, deployment, and ROI

Duration: 1 hour



Data Governance

The Data Imperative

**"Organizations don't have AI problems;
they have data problems that AI exposes."**

Plan for 60–80% of GenAI project time
to be spent on data preparation.

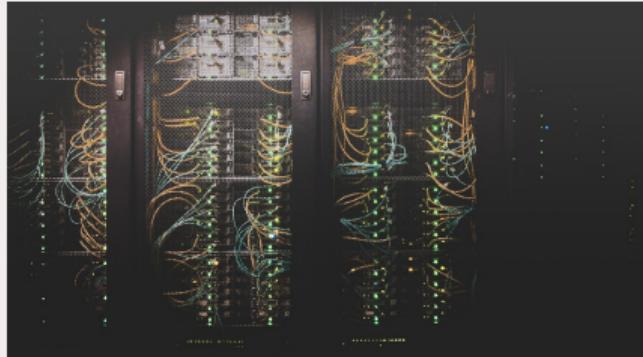
Data Strategy Precedes AI Strategy

The Data Hierarchy of Needs:

- 1 Data Collection** — Foundation
- 2 Clean Data** — Must start here
- 3 Analytics & Reporting**
- 4 AI/ML** — Most start here (mistake)

Reality Check

Fortune 500 expected 4 months for GenAI. Actual: 15 months. Root cause: Data readiness.



Data Requirements for GenAI

- **Training Data:** Building/fine-tuning models
Strategic value: Competitive moat
- **Context Data (RAG):** Grounding model outputs
Strategic value: Accuracy & relevance
- **Operational Data:** Real-time model inputs
Strategic value: Timeliness

Quality Dimensions: Accuracy, Completeness, Consistency, Timeliness, Representativeness

Global Privacy Regulations

- **GDPR** (EU): Up to 4% global revenue
- **CCPA/CPRA** (California):
Per-violation penalties
- **PIPL** (China): Up to 5% revenue
- **LGPD** (Brazil): Up to 2% revenue
- **POPIA** (South Africa): Up to 10M ZAR

Global Trend

Design AI systems with privacy by default.



GenAI-Specific Privacy Concerns

- 1 Training Data Privacy:** Was personal data used with consent?
- 2 Inference Privacy:** Can model be manipulated to reveal data?
- 3 Output Privacy:** Do outputs contain personal information?
- 4 Conversation Privacy:** Who accesses user interactions?
- 5 Derived Data:** Are new personal insights generated?

The Consent Challenge

Traditional consent breaks down: capabilities hard to explain, data use unpredictable, untraining technically difficult.

Data Governance Framework

Key Components

- Data inventory & classification
- Access controls
- Consent management
- Retention policies
- Audit trails

Best Practices

- Minimize data collection
- Purpose limitation
- Regular compliance audits
- Incident response plans
- Cross-border controls

User Rights to Support

- **Right to Access:** Users request all data held about them
- **Right to Erasure:** Users request deletion
- **Right to Portability:** Data in machine-readable format
- **Right to Rectification:** Correct inaccurate data
- **Right to Object:** Object to certain processing
- **Automated Decision Rights:** Human review of AI decisions

China's AI Regulatory Framework

The world's most comprehensive AI regulations:

- **Algorithm Recommendations** (2022): Internet services
- **Deep Synthesis** (2023): Deepfakes, synthetic media
- **GenAI Service Measures** (2023): All public GenAI
- **AIGC Labeling** (Sept 2025): Mandatory AI content labels
- **National Standards** (Nov 2025): Security & governance

Scale: 350+ LLMs filed. 1.57M AI patents (38.6% of global total).



Product Development

The GenAI Development Reality

Key Statistics (2025)

Only **5%** of AI pilots achieve rapid revenue acceleration

67% success rate for purchasing/partnering

22% success rate for internal builds

46% have no structured ROI measurement

GenAI has entered the “Trough of Disillusionment”

Why Traditional Project Management Fails

Traditional

- Fixed requirements
- Binary success
- Predictable timeline
- Deterministic testing

GenAI

- Emergent requirements
- Probabilistic success
- Uncertain timeline
- Statistical testing



Implication

Waterfall always fails. Agile is better but less efficient.

The AI Project Lifecycle

- 1 Problem Framing** (Often Skipped): Should AI solve this?
- 2 Data Assessment:** Inventory, gaps, quality
- 3 Proof of Concept** (4–8 weeks): Time-boxed experimentation
- 4 Pilot:** Limited production, controlled blast radius
- 5 Production & Scale:** Infrastructure, monitoring
- 6 Operations:** Performance monitoring, retraining

Rule of Thumb

Budget for 2–3 PoCs failing for every success.

Phase Gates for GenAI

- **Gate 0:** Business case, feasibility, ethics screening
- **Gate 1:** Requirements, data availability, build vs. buy
- **Gate 2:** Technical validation, benchmarks, user feedback
- **Gate 3:** Production-grade, security & ethics review
- **Gate 4:** Controlled deployment, monitoring setup
- **Gate 5:** Full deployment, continuous improvement

Kill Criteria: Define Before Starting

- **Technical:** Can't achieve accuracy threshold
- **Economic:** Cost exceeds value
- **Timeline:** 6-month delay, no path forward
- **Ethical:** Can't mitigate bias
- **Security:** Can't protect data
- **Regulatory:** Unacceptable compliance risk
- **Strategic:** Market opportunity gone



Implementation Patterns

1 Co-Pilot / Augmentation

AI assists; humans decide. *Best for: High-stakes, building trust*

2 Automation with Exceptions

AI handles routine; humans handle exceptions. *Best for: High-volume*

3 Full Automation

AI autonomous with monitoring. *Best for: Low-stakes, speed critical*

4 Internal Tool

AI assists employees only. *Best for: Building capability, lower risk*

Build vs. Buy Decision

- **Build from Scratch:** \$10M–\$100M+; 12–24 months
Only if. Massive data advantage
- **Fine-Tune:** \$10K–\$1M; weeks to months
Best for. Domain-specific tasks
- **RAG:** \$10K–\$100K; weeks
Best for. Current/proprietary information
- **Prompt Engineering:** \$1K–\$10K; days to weeks
Best for. Quick wins
- **Buy SaaS:** Variable; days
Best for. Non-differentiating capabilities

Success Metrics

Avoid Vanity Metrics:

- ✗ "We deployed an AI model"
- ✗ "95% accuracy" (on what?)

Focus on Business Outcomes:

- ✓ Customer satisfaction improved by X%
- ✓ Time to resolution decreased by Y hours
- ✓ Cost per transaction reduced by \$Z
- ✓ Employee time redirected to higher-value work

Four-Layer Monitoring Framework

- 1 Infrastructure:** Latency, error rates, throughput, cost
- 2 Model Performance:** Accuracy, hallucination rate, drift
- 3 Business:** Adoption, task completion, satisfaction, revenue
- 4 Risk:** Incidents, near-misses, compliance, complaints

Principle

You can't improve what you don't measure. Monitor from day one.

ROI Reality (2025)

- Average ROI: **3.7x** per dollar (IDC/Microsoft)
- Top performers: **\$10.3** return per dollar
- 74% meeting or exceeding expectations (Deloitte)
- **46% have no structured ROI measurement**

Timeline Expectations:

- Chatbots, RPA: 6–12 months
- Operational efficiency: 12–24 months
- Revenue generation: 18–36 months

Total Cost of Ownership

Initial Costs

- Infrastructure (GPUs)
- Software licenses
- Integration
- Data preparation
- Training

Ongoing Costs

- Compute resources
- API fees
- Model maintenance
- Monitoring
- Personnel

Hidden Costs: Compliance, legal/IP, incidents, technical debt, failed pilots

Minimum Viable AI Team

- **Executive Sponsor** (10–20%): Alignment, resources, blockers
- **Product Owner** (Full-time): Requirements, prioritization
- **Data Engineer** (Full-time): Pipelines, quality
- **ML Engineer** (Full-time): Model development
- **Domain Expert** (25–50%): Business logic, validation
- **MLOps Engineer**: Deployment, monitoring

Part 1 Key Takeaways

Summary

- 1 **Data First:** 60–80% of GenAI time is data preparation
- 2 **Privacy by Design:** Global regulations require it
- 3 **Expect Failure:** Budget for 2–3 PoCs failing per success
- 4 **Define Kill Criteria:** Before emotional investment
- 5 **Measure Everything:** Connect to business outcomes
- 6 **Build the Right Team:** Minimum viable AI team

Discussion Questions

- 1** What is the current state of data readiness in your organization?

- 2** Have you defined clear kill criteria for your AI projects?

- 3** How are you measuring ROI on AI investments today?

- 4** Do you have the right team composition for AI success?

Thank You



www.hdx.edu

info@hdx.edu

@HappyDigitalX

Continue to Part 2: Ethics, Security & Imple