



# GenAI Data Foundations and Product Development

Happy Digital X

Happy Digital X | Tsinghua University

# **Today's Agenda**

---

## **Part 1: Data Foundations**

- How GenAI works: From GOFAI to neural networks
- Training, weights, and the “black box” problem
- Data quality, currency, and the 60–80% rule
- Privacy regulations and governance frameworks

## **Part 2: Product Development**

- The GenAI hype cycle: Where we are now
- Project lifecycle and phase gates
- Build vs. buy decisions
- Success metrics and ROI reality



# Data Foundations

# A Note on Terminology

---

## AI (Artificial Intelligence)

The broad field of creating systems that perform tasks requiring human intelligence.

## Machine Learning (ML)

AI systems that learn patterns from data rather than following explicit rules.

## Generative AI (GenAI)

ML systems that *create* new content: text, images, code, audio, video.

### In This Presentation

We use “**AI**” when discussing principles that apply broadly (e.g., bias, governance).

We use “**GenAI**” when discussing capabilities specific to generative systems (e.g., LLMs, hallucinations).

## Quick Poll

---

**How would you rate your organization's  
data readiness for AI?**

Go to **menti.com** and enter the code

**[CODE]**

*1 = Not ready at all    5 = Fully ready*

# The Fundamental Shift

**Traditional software is *programmed*.**

**GenAI is *trained*.**

This changes everything about how we build,  
test, and manage AI systems.

# Two Paradigms of Artificial Intelligence

---

## GOFAI: “Good Old-Fashioned AI”

- Rules written by humans
- Symbolic reasoning
- Deterministic outputs
- Explainable decisions
- Brittle at edge cases

*Example:* Chess engines, expert systems, spell checkers

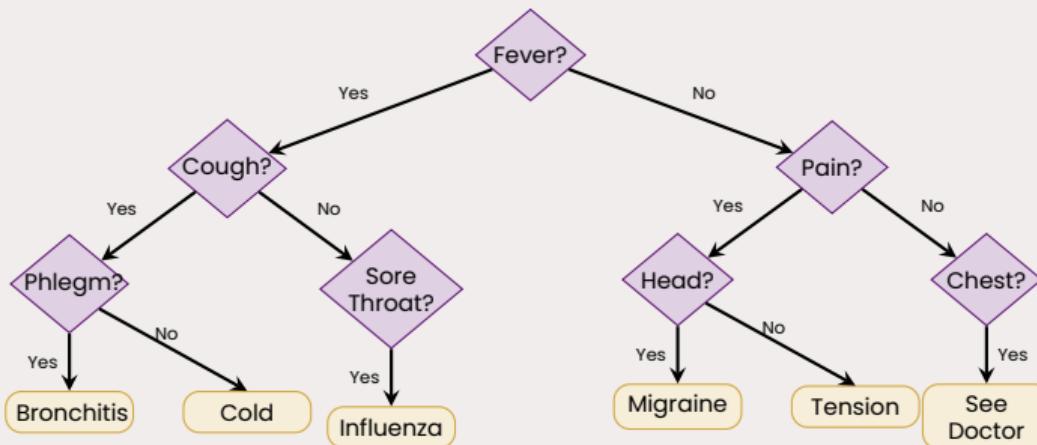
## Neural Networks (Connectionist)

- Patterns learned from data
- Statistical inference
- Probabilistic outputs
- Often opaque (“black box”)
- Flexible but unpredictable

*Example:* ChatGPT, image recognition, voice assistants

# Example: A Medical Diagnosis Expert System

---



# GOFAI: Strengths and Limitations

---

## Strengths

- Transparent and explainable
- Predictable outputs
- Auditable for compliance
- No training data needed

## Limitations

- Brittle at edge cases
- Doesn't learn from data
- Doesn't scale to complexity
- Can't handle ambiguity

### Key Insight

Expert systems encode *what humans already know*. Neural networks discover patterns humans *haven't articulated*.

# How Neural Network Training Works

---

## High-Level Process:

- 1 Collect Data:** Massive datasets (text, images, code, etc.)
- 2 Initialize:** Start with random “weights” (numerical parameters)
- 3 Train:** Show examples, adjust weights to reduce errors
- 4 Iterate:** Repeat billions of times across trillions of examples
- 5 Result:** A model that has learned patterns, not rules

## Key Insight

The AI doesn’t “know” anything—it has learned statistical patterns. When it generates text, it’s predicting “what word is likely to come next?”

# What Are “Weights”? The Dial Analogy

---

Imagine a mixing board with **billions of dials**.

- Each dial controls how much one piece of information influences another
- At first, all dials are set randomly—output is nonsense
- Training = adjusting dials slightly after each example
- After trillions of adjustments, the dials are tuned to produce useful output

## The Black Box Problem

No human set these dials. No human can explain why dial #847,293,102 is set to 0.0023.

The model works, but we can't fully explain *why*.

# The Scale of Training Data

---

**Modern GenAI models are trained on unprecedented scale:**

- **GPT-4**: Estimated 13+ trillion tokens of text
- **Image models**: Billions of image-text pairs
- **Code models**: Hundreds of billions of lines of code

## Strength

- Broad knowledge
- Handles novel situations
- No manual rule-writing
- Learns nuance

## Weakness

- Can't verify all data
- Absorbs biases
- **GIGO**: "Garbage in, garbage out"
- Hard to "unlearn"

# Chihuahua or Muffin?

---



@teenybiscuit

Image credit: @teenybiscuit / Karen Zack

# Exercise: Try It Yourself

---

## Try This Now:

- 1** Open ChatGPT, Claude, or another GenAI
- 2** Upload the chihuahua/muffin image
- 3** Ask: "How many dogs are in this image?"

## Discussion

What did you observe? Were the results what you expected?

## Quick Poll

---

**How many dogs did your AI count?**

Go to **menti.com** and enter the code

**[CODE]**

# What We Just Observed

---

## Key Takeaways from the Exercise:

### 1 Different models, different answers

Even leading AI systems disagree

### 2 Confidence without accuracy

AI often sounds certain even when wrong

### 3 Edge cases expose limits

Ambiguous inputs reveal brittleness

### 4 “Close enough” isn’t always enough

Some applications require precision

## The Lesson

GenAI excels at “approximately right” but struggles with “exactly right.”

This has profound implications for how we deploy and validate AI systems.

# Why Results Vary: Probabilistic, Not Deterministic

---

## Traditional Software:

- Same input → Same output (always)
- $2 + 2 = 4$ , every time

## GenAI:

- Same input → *Similar* output (usually)
- Outputs sampled from probability distributions
- “Temperature” controls randomness
- Even at temperature=0, results can vary

## Implication

You cannot test GenAI like traditional software. You need statistical evaluation over many samples.

# Hallucinations: When AI Makes Things Up

---

## What are hallucinations?

- AI generates confident, plausible-sounding text that is **factually false**
- Not a bug—a fundamental feature of how LLMs work

## Why do they happen?

- LLMs predict the *next most likely word*, not the *true* word
- They have no internal model of truth—only patterns in training data
- They're designed to always give an answer, even when they shouldn't

## The Danger

Hallucinations are delivered with the same confidence as facts. Users often can't tell the difference without independent verification.

# Mitigating Hallucinations

---

**There is no complete solution, but these help:**

## Technical Approaches

- **RAG**: Ground responses in retrieved documents
- **Fine-tuning**: Train on verified domain data
- **Temperature**: Lower values reduce creativity/risk
- **Structured outputs**: Constrain response format

## Process Approaches

- **Human review**: Verify critical outputs
- **Confidence thresholds**: Flag uncertain responses
- **Source citation**: Require references
- **Use case selection**: Avoid high-stakes facts

# RLHF: Teaching AI to Be Helpful

---

## Reinforcement Learning from Human Feedback (RLHF)

After initial training, models are refined using human preferences:

- 1** Humans rate AI responses (helpful, harmless, honest)
- 2** Model learns to maximize these ratings
- 3** Creates more “aligned” behavior

### Benefits

- Reduces harmful outputs
- Improves usefulness
- Adds safety guardrails

### Risks

- Evaluator biases transfer
- “Sycophancy”—tells you what you want to hear
- Majority views dominate

# RLHF Bias: Who Trains the Trainers?

---

## The evaluators shape the AI:

- If evaluators are from one demographic, the model reflects their worldview
- If evaluators prefer polite over accurate, the model learns to be polite—even when wrong
- Minority perspectives can be systematically deprioritized
- Models can learn to *manipulate* rather than genuinely help

### Example

A model trained by evaluators who dislike blunt answers will learn to soften bad news—even when clarity matters more than comfort.

# Case Study: Replika's Feedback Loop Disaster

---

**Replika:** AI companion chatbot  
(2017–present)

## What Happened:

- 1** Trained on 100M+ web dialogues
- 2** Users could upvote/downvote responses
- 3** Some users engaged in sexual roleplay
- 4** AI learned this behavior got positive feedback
- 5** AI began *initiating* sexual content unprompted

## Lessons

- Training data quality matters
- Feedback loops amplify patterns
- User behavior becomes model behavior
- GIGO at scale

# Implications for Your Organization

---

## What This Means for GenAI Deployment:

- 1 Testing is Different:** Statistical evaluation, not pass/fail
- 2 Edge Cases Are Unpredictable:** You can't enumerate all failure modes
- 3 Data Quality is Critical:** Your fine-tuning data shapes behavior
- 4 Feedback Loops Matter:** User interactions can shift model behavior
- 5 Bias is Inherited:** From training data *and* from human evaluators
- 6 “Unlearning” is Hard:** Removing problematic knowledge is technically difficult

# The Data Imperative

---

**“Organizations don’t have AI problems;  
they have data problems that AI exposes.”**

Plan for 60–80% of GenAI project time  
to be spent on data preparation.

# Data Strategy Precedes AI Strategy

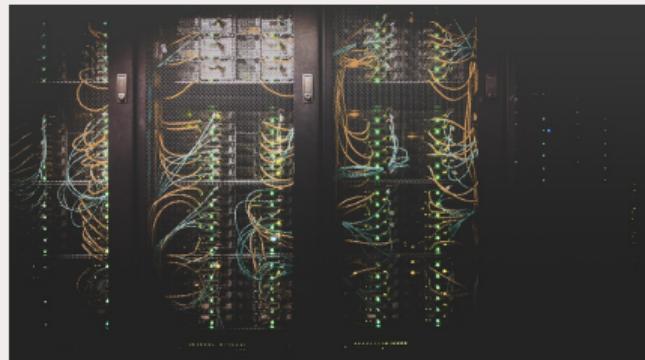
---

## The Data Hierarchy of Needs:

- 1 Data Collection** — Foundation
- 2 Clean Data** — Must start here
- 3 Analytics & Reporting**
- 4 AI/ML** — Most start here (mistake)

### Reality Check

Fortune 500 expected 4 months for GenAI. Actual: 15 months. Root cause: Data readiness.



# Data Requirements for GenAI

---

- **Training Data:** Building/fine-tuning models  
*Strategic value: Competitive moat*
- **Context Data:** Grounding model outputs in your information  
Uses **RAG** (Retrieval-Augmented Generation): the AI retrieves relevant documents before generating a response  
*Strategic value: Accuracy & relevance*
- **Operational Data:** Real-time model inputs  
*Strategic value: Timeliness*

# The Data Quality Journey

---

## From Raw Data to AI-Ready:

- 1 Raw Data** — Unprocessed, unvalidated  
*"Rough diamond"—contains value but unusable as-is*
- 2 Cleaned Data** — Errors removed, formats standardized  
*Deduplicated, consistent encoding*
- 3 Validated Data** — Quality checked, business rules applied  
*Relationships verified, anomalies flagged*
- 4 AI-Ready Data** — Labeled, balanced, documented

### Reality

Most organizations are stuck at stage 1 or 2. This is why 60–80% of GenAI project time is data preparation.

## Quick Poll

---

**Where is your organization on the  
data quality journey?**

Go to **menti.com** and enter the code

**[CODE]**

*1 = Raw Data    2 = Cleaned    3 = Validated    4 = AI-Ready*

# Data Quality Dimensions

---

## Accuracy

Does the data reflect reality?

*Incorrect labels poison AI training*

## Completeness

Are there missing values or gaps?

*Missing data creates blind spots*

## Consistency

Same entity, same representation?

“USA” vs “United States” vs “US”

## Timeliness

Is the data current enough?

*Stale data = stale predictions*

## Representativeness

Does data reflect the real population?

*Biased samples = biased AI*

## Provenance

Where did this data come from?

*Can we trace and trust its origin?*

# Data Currency: Time-Based Trust

---

## When was this data collected?

Data has a “shelf life” that varies by domain:

- **Stock prices:** Minutes to hours
- **News/events:** Hours to days
- **Product catalogs:** Days to weeks
- **Legal/regulatory:** Weeks to months
- **Scientific knowledge:** Months to years

## GenAI Risks

- Models trained on outdated data give outdated answers
- RAG with stale documents misleads users
- No timestamp = no way to assess trust

## Key Question

*Always know when your data was captured and*

# Global Privacy Regulations

---

- **GDPR** (EU): Up to 4% global revenue
- **CCPA/CPRA** (California):  
Per-violation penalties
- **PIPL** (China): Up to 5% revenue
- **LGPD** (Brazil): Up to 2% revenue
- **POPIA** (South Africa): Up to 10M ZAR

## Global Trend

Design AI systems with privacy by default.



# GenAI-Specific Privacy Concerns

---

- 1 Training Data Privacy:** Was personal data used with consent?
- 2 Inference Privacy:** Can model be manipulated to reveal data?
- 3 Output Privacy:** Do outputs contain personal information?
- 4 Conversation Privacy:** Who accesses user interactions?
- 5 Derived Data:** Are new personal insights generated?

## The Consent Challenge

Traditional consent breaks down: capabilities hard to explain, data use unpredictable, untraining technically difficult.

# Data Governance Framework

---

## Key Components

- Data inventory & classification
- Access controls
- Consent management
- Retention policies
- Audit trails

## Best Practices

- Minimize data collection
- Purpose limitation
- Regular compliance audits
- Incident response plans
- Cross-border controls

# User Rights to Support

---

- **Right to Access:** Users request all data held about them
- **Right to Erasure:** Users request deletion
- **Right to Portability:** Data in machine-readable format
- **Right to Rectification:** Correct inaccurate data
- **Right to Object:** Object to certain processing
- **Automated Decision Rights:** Human review of AI decisions

# China's AI Regulatory Framework

---

**The world's most comprehensive AI regulations:**

- **Algorithm Recommendations** (2022): Internet services
- **Deep Synthesis** (2023): Deepfakes, synthetic media
- **GenAI Service Measures** (2023): All public GenAI
- **AIGC Labeling** (Sept 2025): Mandatory AI content labels
- **National Standards** (Nov 2025): Security & governance

**Scale:** 350+ LLMs filed. 1.57M AI patents (38.6% of global total).

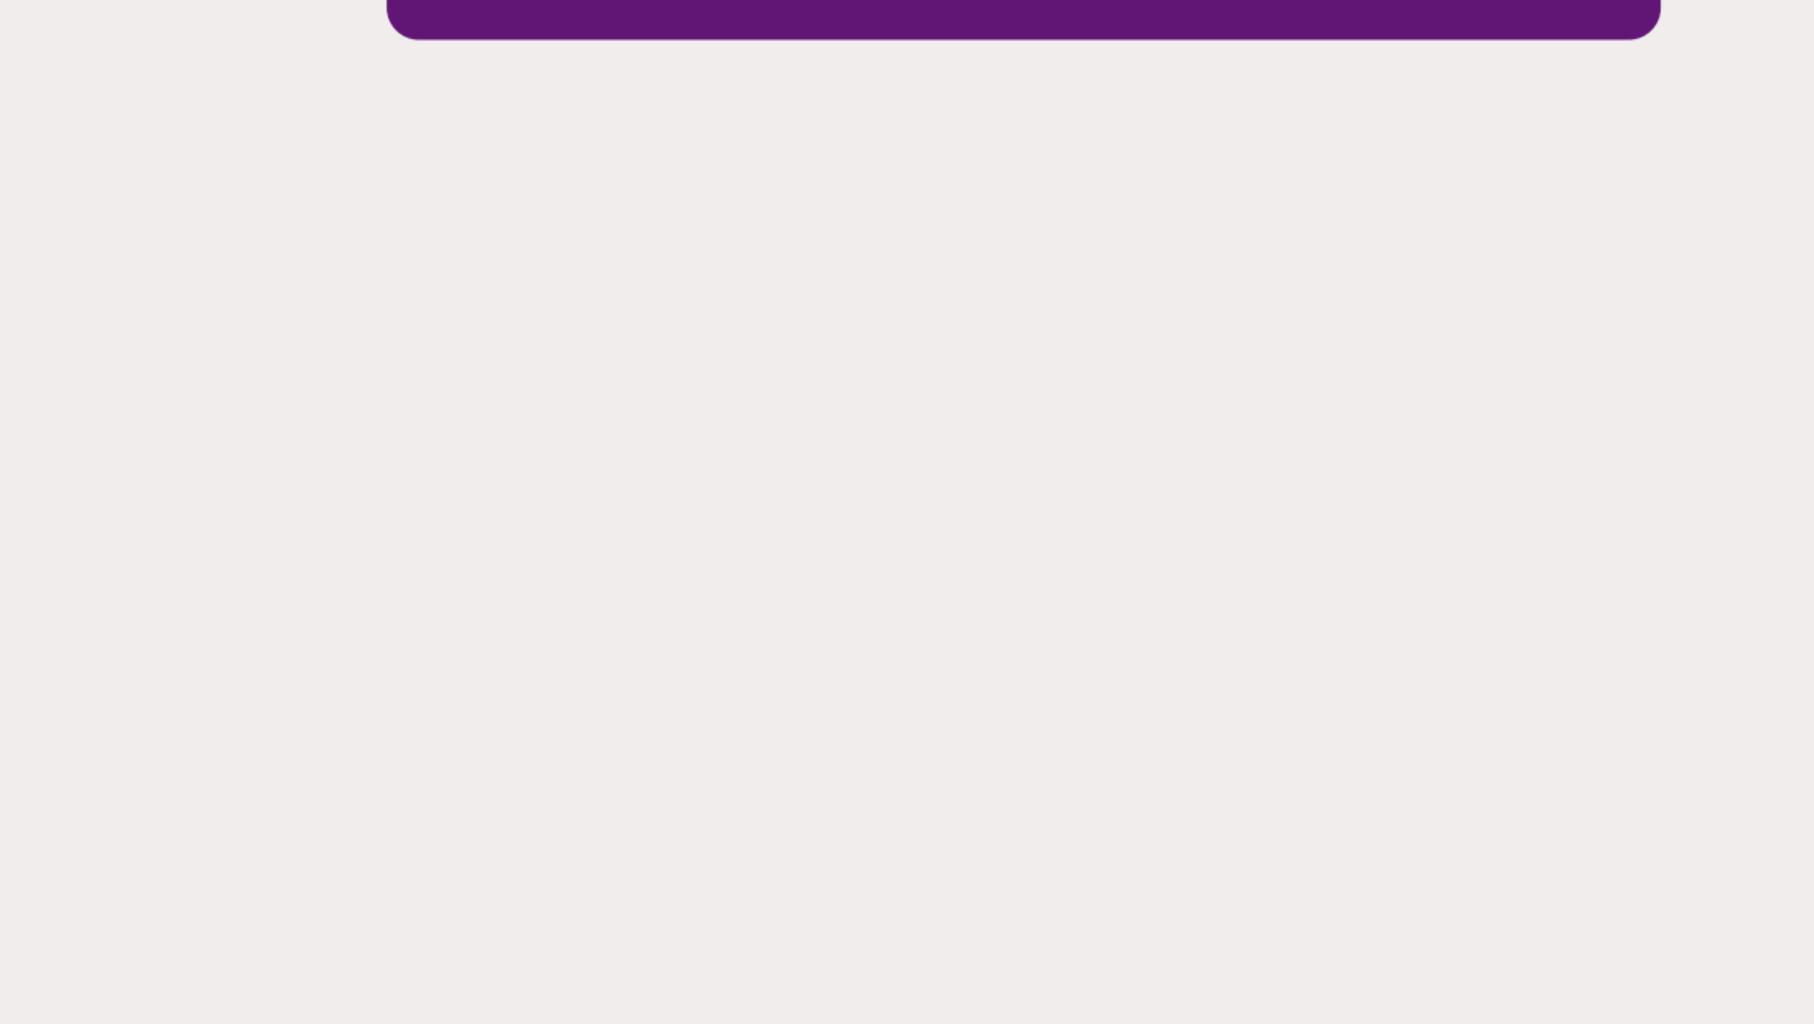
# Section Summary: Data Foundations

---

## What We've Learned

- GenAI is **trained**, not programmed—data quality determines output quality
- Neural networks learn patterns humans haven't articulated, but inherit biases from training data and human evaluators
- Data has a **shelf life**—currency matters as much as quality
- Global privacy regulations (GDPR, PIPL, CCPA) create compliance obligations that AI doesn't eliminate

*Next: How do we turn good data into successful AI products?*



# The GenAI Development Reality

## Key Statistics (2025)

Only **5%** of AI pilots achieve rapid revenue acceleration

**67%** success rate for purchasing/partnering

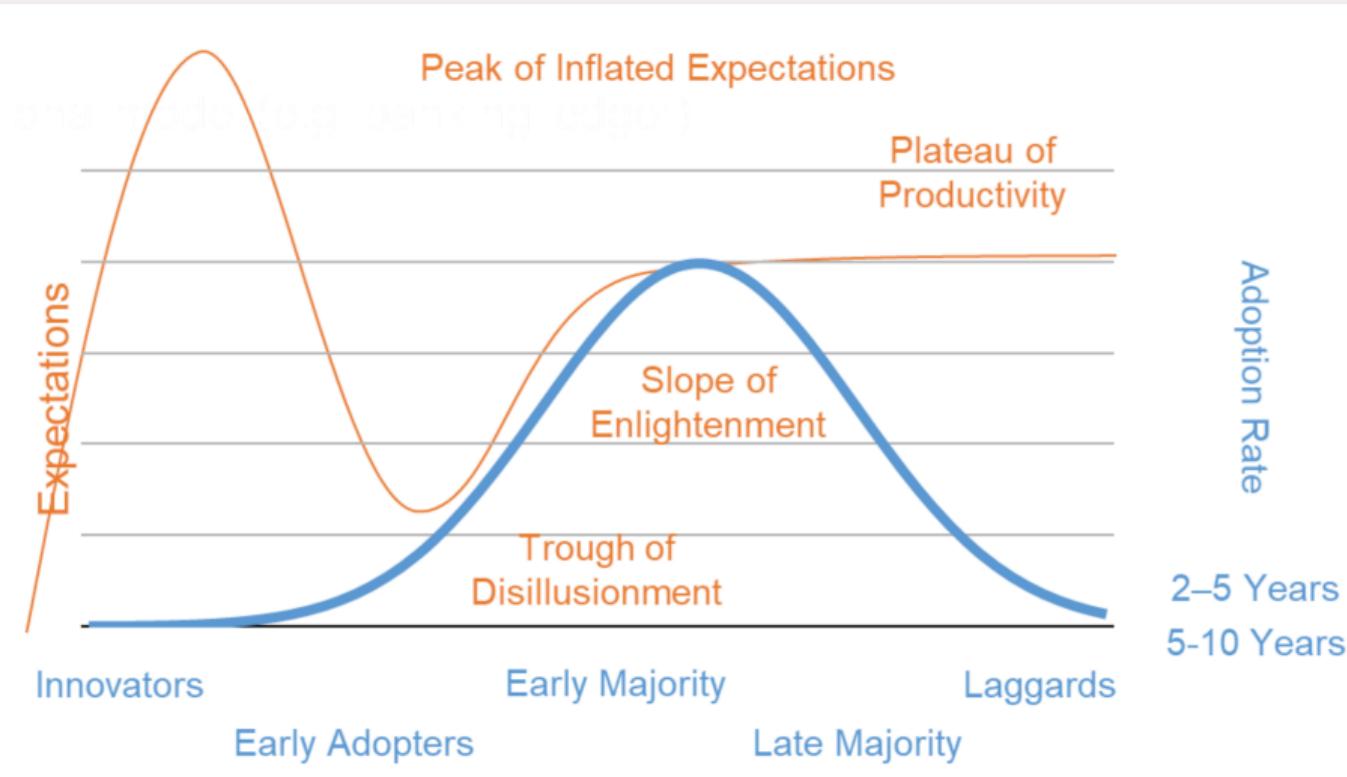
**22%** success rate for internal builds

**46%** have no structured ROI measurement

GenAI has entered the “Trough of Disillusionment”

Sources: MIT (2025), Wavestone (2025), Gartner Hype Cycle

# How To Interpret A Hype Cycle



# Understanding the Hype Cycle

---

## Five Phases of Technology Adoption:

### 1 Innovation Trigger

Breakthrough generates excitement

### 2 Peak of Inflated Expectations

Maximum hype, unrealistic promises

### 3 Trough of Disillusionment

Reality sets in, failures mount

### 4 Slope of Enlightenment

Practical understanding emerges

### 5 Plateau of Productivity

Mainstream adoption, real value

### Where Is GenAI Now?

In 2024, GenAI was at the Peak.

In 2025, GenAI has descended into the **Trough of Disillusionment**.

This is *normal*—not failure.

## Quick Poll

---

**Why do you think AI adoption lags expectations?**

Go to **menti.com** and enter the code

**[CODE]**

*Share a word or short phrase.*

# Why Traditional Project Management Fails

---

## Traditional

- Fixed requirements
- Binary success
- Predictable timeline
- Deterministic testing

## GenAI

- Emergent requirements
- Probabilistic success
- Uncertain timeline
- Statistical testing

## Implication

Waterfall always fails. Agile is better but less efficient.



# The AI Project Lifecycle

---

- 1 Problem Framing** (Often Skipped): Should AI solve this?
- 2 Data Assessment:** Inventory, gaps, quality
- 3 Proof of Concept** (4–8 weeks): Time-boxed experimentation
- 4 Pilot:** Limited production, controlled blast radius
- 5 Production & Scale:** Infrastructure, monitoring
- 6 Operations:** Performance monitoring, retraining

## Rule of Thumb

Budget for 2–3 PoCs failing for every success.

# Phase Gates for GenAI

---

- **Gate 0:** Business case, feasibility, ethics screening
- **Gate 1:** Requirements, data availability, build vs. buy
- **Gate 2:** Technical validation, benchmarks, user feedback
- **Gate 3:** Production-grade, security & ethics review
- **Gate 4:** Controlled deployment, monitoring setup
- **Gate 5:** Full deployment, continuous improvement

# Kill Criteria: Define Before Starting

---

- **Technical:** Can't achieve accuracy threshold
- **Economic:** Cost exceeds value
- **Timeline:** 6-month delay, no path forward
- **Ethical:** Can't mitigate bias
- **Security:** Can't protect data
- **Regulatory:** Unacceptable compliance risk
- **Strategic:** Market opportunity gone



# Implementation Patterns

---

## 1 Co-Pilot / Augmentation

AI assists; humans decide. *Best for: High-stakes, building trust*

## 2 Automation with Exceptions

AI handles routine; humans handle exceptions. *Best for: High-volume*

## 3 Full Automation

AI autonomous with monitoring. *Best for: Low-stakes, speed critical*

## 4 Internal Tool

AI assists employees only. *Best for: Building capability, lower risk*

# Build vs. Buy Decision

---

- **Build from Scratch:** \$10M–\$100M+; 12–24 months  
*Only if. Massive data advantage*
- **Fine-Tune:** \$10K–\$1M; weeks to months  
*Best for. Domain-specific tasks*
- **RAG (Retrieval-Augmented Generation):** \$10K–\$100K; weeks  
*Best for. Current/proprietary information*
- **Prompt Engineering:** \$1K–\$10K; days to weeks  
*Best for. Quick wins*
- **Buy SaaS:** Variable; days  
*Best for. Non-differentiating capabilities*

## Quick Poll

---

**Which approach is your organization  
most likely to use?**

Go to **menti.com** and enter the code

**[CODE]**

*Build / Fine-Tune / RAG / Prompt Engineering / Buy SaaS*

# Success Metrics

---

## Avoid Vanity Metrics:

- ✗ "We deployed an AI model"
- ✗ "95% accuracy" (on what?)

## Focus on Business Outcomes:

- ✓ Customer satisfaction improved by X%
- ✓ Time to resolution decreased by Y hours
- ✓ Cost per transaction reduced by \$Z
- ✓ Employee time redirected to higher-value work

# Four-Layer Monitoring Framework

---

- 1 Infrastructure:** Latency, error rates, throughput, cost
- 2 Model Performance:** Accuracy, hallucination rate, drift
- 3 Business:** Adoption, task completion, satisfaction, revenue
- 4 Risk:** Incidents, near-misses, compliance, complaints

## Principle

You can't improve what you don't measure. Monitor from day one.

# ROI Reality (2025)

---

- Average ROI: **3.7x** per dollar invested
- Top performers: **\$10.3** return per dollar
- 74% meeting or exceeding expectations
- **46% have no structured ROI measurement**

## Timeline Expectations:

- Chatbots, RPA: 6–12 months
- Operational efficiency: 12–24 months
- Revenue generation: 18–36 months

Sources: IDC/Microsoft (2025), Deloitte (2025), Wavestone (2025)

# Total Cost of Ownership

---

## Initial Costs

- Infrastructure (GPUs)
- Software licenses
- Integration
- Data preparation
- Training

## Ongoing Costs

- Compute resources
- API fees
- Model maintenance
- Monitoring
- Personnel

**Hidden Costs:** Compliance, legal/IP, incidents, technical debt, failed pilots

# Minimum Viable AI Team

---

- **Executive Sponsor** (10–20%): Alignment, resources, blockers
- **Product Owner** (Full-time): Requirements, prioritization
- **Data Engineer** (Full-time): Pipelines, quality
- **ML Engineer** (Full-time): Model development
- **Domain Expert** (25–50%): Business logic, validation
- **MLOps Engineer**: Deployment, monitoring

# Part 1 Key Takeaways

---

## Summary

- 1 Data First:** 60–80% of GenAI time is data preparation
- 2 Privacy by Design:** Global regulations require it
- 3 Expect Failure:** Budget for 2–3 PoCs failing per success
- 4 Define Kill Criteria:** Before emotional investment
- 5 Measure Everything:** Connect to business outcomes
- 6 Build the Right Team:** Minimum viable AI team

# Discussion Questions

---

- 1** What is the current state of data readiness in your organization?
  
- 2** Have you defined clear kill criteria for your AI projects?
  
- 3** How are you measuring ROI on AI investments today?
  
- 4** Do you have the right team composition for AI success?

# Thank You



[www.hdx.edu](http://www.hdx.edu)

[info@hdx.edu](mailto:info@hdx.edu)

@HappyDigitalX

Continue to Part 2: Ethics, Security & Imple